Volatility Arbitrage via the Variance Risk Premium:

Trading the IV–Forecast RV Spread on SPX Options

Group 4
MS&E 244: Statistical Arbitrage
Stanford University

Winter 2026

1 Introduction

1.1 Motivation and Trading Style

Options markets embed forward-looking expectations about future volatility through implied volatility (IV). A large body of empirical research documents that risk-neutral implied variance systematically exceeds the variance that subsequently realizes under the physical measure—a phenomenon known as the *variance risk premium* (VRP). This gap compensates option sellers for bearing crash risk and reflects investors' aggregate demand for downside protection [Bollerslev et al., 2009, Carr and Wu, 2009].

Our strategy exploits this premium systematically. We trade **delta-hedged ATM straddles on SPX options** at daily frequency, entering positions when the spread between IV and our forecast of realized volatility (RV) is statistically extreme. This isolates pure volatility exposure: our PnL depends on whether implied vol over- or under-estimates future realized vol, not on the direction of the equity market.

1.2 Literature Review

The variance risk premium has been studied extensively. Carr and Madan [1998] and Demeterfi et al. [1999] develop the theoretical framework for variance swaps and model-free implied variance. Bollerslev et al. [2009] show that the VRP predicts equity returns and Carr and Wu [2009] decompose the variance premium across currencies and equities. Bondarenko [2003] documents the systematic overpricing of put options, consistent with a negative variance risk premium from the seller's perspective.

On the realized volatility side, Andersen et al. [2003] pioneer high-frequency RV estimation. Corsi [2009] introduces the Heterogeneous Autoregressive (HAR-RV) model, which captures the multi-scale persistence of volatility using daily, weekly, and monthly lag components. Hansen et al. [2012] extend GARCH to incorporate realized measures. For risk-neutral distribution extraction, Breeden and Litzenberger [1978] provide the foundational result linking option prices to state-contingent claims, and Bakshi et al. [2003] study the link between skewness and option pricing. Gatheral [2006] offers a practitioner-oriented treatment of volatility surface dynamics.

1.3 Strategy Overview

We implement and compare three signal families derived from the VRP literature:

1. **VRP Core Signal**: $S_t = \mathrm{IV}_t - \widehat{\mathrm{RV}}_t$, z-score normalized, using ATM IV and composite RV forecasts from HAR-RV, GARCH, and GJR-GARCH models.
2. **Variance Swap Signal**: Replaces ATM IV with model-free implied variance (MFIV) computed via the discrete Carr–Madan approximation.
3. **Distribution-Based Signal**: Compares the risk-neutral return distribution (extracted from the option surface) to the physically forecast distribution, trading on tail divergence.

On $1M notional capital over 2011–2023, the core VRP strategy (with a 25% stop-loss per trade) generates 93 trades with a Sharpe ratio of 5.26, a 50.5% win rate, and $7.63M cumulative PnL. The strategy is profitable out-of-sample and across most stress periods; a per-trade stop-loss (25% of trade value) limits tail losses so that maximum drawdown is 79.4% rather than the 98%+ that would occur without it (e.g., the February 2018 "Volmageddon" event [Augustin et al., 2021]).

2 Model

2.1 Signal Construction

2.1.1 Implied Volatility Measures

We compute two IV measures for each trading date:

ATM Implied Volatility. We select options with $|\Delta| \in [0.40, 0.60]$ (near-ATM) and compute a delta-weighted average IV for each available expiry. We then linearly interpolate across expiries to obtain a constant-maturity 30-day ATM IV:

$$\mathrm{IV}_{30}^{\mathrm{ATM}} = \mathrm{IV}_{\mathrm{short}} \cdot w + \mathrm{IV}_{\mathrm{long}} \cdot (1 - w), \quad w = \frac{T_{\mathrm{long}} - 30}{T_{\mathrm{long}} - T_{\mathrm{short}}} \tag{1}$$

where T_{short} and T_{long} are the days-to-expiration that bracket the 30-day target.

Model-Free Implied Variance (MFIV). Following Carr and Madan [1998] and Demeterfi et al. [1999], we construct a synthetic variance swap rate by integrating over out-of-the-money (OTM) option prices:

$$\sigma^2_{\mathrm{MFIV}} = \frac{2}{T} \sum_i \frac{\Delta K_i}{K_i^2} e^{rT} Q(K_i) \tag{2}$$

where $Q(K_i)$ is the OTM option mid-price at strike K_i, using puts below the forward and calls above. We restrict strikes to within ±5% of spot and interpolate to 30-day constant maturity.

2.1.2 Realized Volatility Forecasting

We forecast the 22-day-ahead (one-month) realized volatility using an ensemble of three models, each estimated on a rolling 252-day window:

HAR-RV [Corsi, 2009]:

$$\mathrm{RV}_{t+22} = \beta_0 + \beta_1 \,\mathrm{RV}_t^{(1)} + \beta_2 \,\mathrm{RV}_t^{(5)} + \beta_3 \,\mathrm{RV}_t^{(22)} + \varepsilon_t \tag{3}$$

where $\mathrm{RV}_t^{(h)}$ is the average daily realized variance over the trailing h days.

GARCH(1,1) and GJR-GARCH(1,1): We fit standard GARCH and the asymmetric GJR variant to daily log returns, extracting the conditional variance forecast at the 22-day horizon via the `arch` Python library.

Composite Forecast: The final RV forecast is an equal-weight average of all active model forecasts on each date. If the median RV forecast is in variance units (much larger than IV^2), we take the square root to convert to volatility units before computing the spread.

2.1.3 Trading Signals

VRP Signal. The core signal is the IV–RV spread, normalized to a z-score:

$$\mathrm{VRP}_t = \mathrm{IV}_t - \widehat{\mathrm{RV}}_t, \qquad z_t = \frac{\mathrm{VRP}_t - \bar{\mu}_{252}}{\bar{\sigma}_{252}} \tag{4}$$

where $\bar{\mu}_{252}$ and $\bar{\sigma}_{252}$ are the rolling 252-day mean and standard deviation. We generate a discrete signal:

$$\mathrm{signal}_t = \begin{cases} +1 & \text{if } z_t > 1.0 \quad \text{(short vol: IV is rich)} \\ -1 & \text{if } z_t < -1.0 \quad \text{(long vol: IV is cheap)} \\ 0 & \text{otherwise} \end{cases}$$

Variance Swap Signal. Identical z-score framework, substituting $\mathrm{MFIV}_{30}^{1/2}$ for ATM IV.

Distribution Signal. We bin standardized historical returns into 10 quantile buckets and compute the physical probability per bin from trailing 252-day data. Risk-neutral probabilities are derived by rescaling the standard normal CDF by the ratio of implied to realized volatility. The signal measures the signed tail divergence between risk-neutral and physical distributions:

$$D_t = \mathrm{sgn}\left(\sum_{j \in \text{tails}} P_j^{\mathrm{RN}} - P_j^{\mathrm{phys}} \right) \cdot \sum_{j \in \text{tails}} |P_j^{\mathrm{RN}} - P_j^{\mathrm{phys}}| \tag{5}$$

This is then z-score normalized with the same framework as the VRP signal.

2.2 Trade Execution

Instrument. We trade delta-hedged ATM straddles (call + put at $K = S$) on the SPX, with a 22-trading-day holding period and the SPX option multiplier of 100.

Position Sizing. For each trade, we compute the number of straddle contracts as:

$$n = \left\lfloor \frac{\text{Notional}}{\text{Straddle Price} \times 100} \right\rfloor \tag{6}$$

with \$1,000,000 notional per trade. All PnLs are scaled by $n \times 100$.

Delta Hedging. We simulate daily delta rebalancing. The straddle delta is $\Delta_{\text{straddle}} = 2\Phi(d_1) - 1$, which starts near zero for ATM and drifts as the underlying moves. We track the daily PnL from the hedge position rebalancing.

PnL Decomposition.

$$\text{Net PnL} = \underbrace{n \cdot 100 \cdot (\text{Premium} - \text{Payoff})}_{\text{Option PnL}} + \underbrace{n \cdot 100 \cdot \text{Hedge PnL}}_{\text{Delta Hedge}} - \underbrace{\text{Txn Cost}}_{\text{Bid/Ask}} \tag{7}$$

Transaction costs are modeled at 5 bps one-way on the full notional deployed, applied at entry and exit.

Stop-Loss (25% of Trade Value). We exit a trade early when unrealized loss reaches 25% of *that trade's value* (premium deployed), not 25% of portfolio notional. Unrealized PnL is computed by marking the straddle to market using the *realized volatility over the trade so far* (so that when volatility spikes, the mark reflects it and the stop triggers). On stop-out, the realized loss is *capped* at 25% of trade value so that the rule actually limits drawdown in the backtest. Without this, marking with entry IV would understate the loss on vol-spike days and the stop would never fire; with the cap, tail events (e.g., Volmageddon) are contained.

2.3 Implementation

The entire pipeline is implemented in Python with a modular architecture:

- `src/config.py` — Centralized configuration (paths, parameters, FOMC dates)
- `src/data_loader.py` — Options/yield data loading, cleaning, SPX spot derivation
- `src/feature_engineering.py` — ATM IV, MFIV, realized variance, FOMC flags
- `src/rv_models.py` — HAR-RV, GARCH, EGARCH, GJR rolling forecasts
- `src/signals.py` — VRP, skew, term-structure, and distribution signals
- `src/backtest.py` — Delta-hedged straddle simulator with position sizing
- `src/performance.py` — Sharpe, Sortino, drawdown, bootstrap, robustness
- `src/visualization.py` — All plotting routines

All code uses rolling estimation windows and strictly avoids look-ahead bias: models are fit only on data available at each decision date.

3 Methodology

3.1 Data

Options Data. We use daily SPX weekly options from OptionMetrics, containing 5.33 million rows spanning October 2005 through February 2023. Each record includes date, expiration, call/put flag, strike price ($\times 1000$), bid/ask quotes, implied volatility, Greeks ($\Delta, \Gamma, \mathcal{V}, \Theta$), volume, open interest, and forward price.

Preprocessing. We apply the following filters, removing 2.31 million rows (43%):

- Bid $\geq$ \$0.05 (eliminate stale/zero quotes)
- Bid-ask spread $< 100\%$ of mid-price
- Days-to-expiration $\in [7, 60]$
- Moneyness $\in [0.80, 1.20]$ (within $\pm 20\%$ of spot)
- Non-missing implied volatility and delta

After cleaning, 3.02 million option records remain across 3,785 trading dates.

SPX Spot Prices. The dataset does not include a direct spot price column. We derive the daily SPX level from the strike prices of the most at-the-money options (those with $|\Delta|$ closest to 0.50), yielding 3,784 trading days.

Risk-Free Rates. We use the US Treasury yield curve panel covering 15,989 days with maturities from 1 to 360 months. The 1-month yield serves as the short-term risk-free rate.

Event Calendar. We flag 236 FOMC meeting dates from 1996–2025. A ± 7 calendar-day window around each meeting defines the "FOMC window," capturing the elevated implied volatility regime surrounding monetary policy announcements. Of our 3,763 feature-table dates, 1,312 (35%) fall within an FOMC window.

3.2 Strategy Testing Scheme

Full-Sample Backtest. The primary backtest runs over the full signal period (April 2011 – February 2023; 2,782 signal days, 93 trades with 25% stop-loss). All RV model estimation uses strictly rolling 252-day windows, so the first year of data serves as a burn-in.

Out-of-Sample Test. We split the signal chronologically: the first 60% of dates form the *in-sample* training period (April 2011 – September 2018), and the remaining 40% form the *out-of-sample* test period (September 2018 – February 2023). The signal construction (z-score rolling window) uses only past data in both periods, so the split tests generalization of the signal's statistical properties.

Walk-Forward Analysis. We divide the full period into 5 non-overlapping windows of approximately equal length and backtest each independently.

Trading Costs. Transaction costs are set at 5 basis points one-way on the full straddle notional, applied at both entry and exit. We test sensitivity from 0 to 20 bps.

No Look-Ahead Bias. Every computation (rolling RV, HAR/GARCH fitting, z-score normalization) uses only data available at time t. The option data is date-stamped and matched exactly. Position entry occurs on the signal date; the price path used for delta hedging starts from the next day's close.

4 Results

4.1 Full-Sample Performance

Table 1 summarizes the baseline VRP strategy performance over the full period.

Table 1: Full-sample performance of the VRP (ATM IV) strategy with 25% stop-loss ($1M notional, 2011–2023).

Metric	**Value**
Number of Trades	93
Trades/Year	8.0
Win Rate	50.5%
Total Net PnL	$7,625,207
Average Trade PnL	$81,991
Median Trade PnL	$1,939
Profit Factor	1.87
Annualized Return	165.2%
Annualized Volatility	31.4%
Sharpe Ratio	5.26
Sortino Ratio	6.46
Calmar Ratio	2.08
Maximum Drawdown	−79.4%
Return Skewness	−1.19
Return Kurtosis	11.7
VaR (95%)	−1.56% daily
CVaR (95%)	−3.56% daily

Benchmark (buy-and-hold SPX): Sharpe 0.62, ann. return 10.9%, max drawdown −33.0%. Strategy alpha (ann.) 88.5%, information ratio 2.44. Capacity (1% of SPX options ADV): $500M.

The strategy generates strong risk-adjusted returns (Sharpe 5.26) with a 25% stop-loss limiting per-trade losses: the average winner ($348K) exceeds the average loser ($190K). Negative skewness (−1.19) and elevated kurtosis (11.7) reflect some left-tail risk, but maximum drawdown is 79.4% rather than 98%+ without the stop.

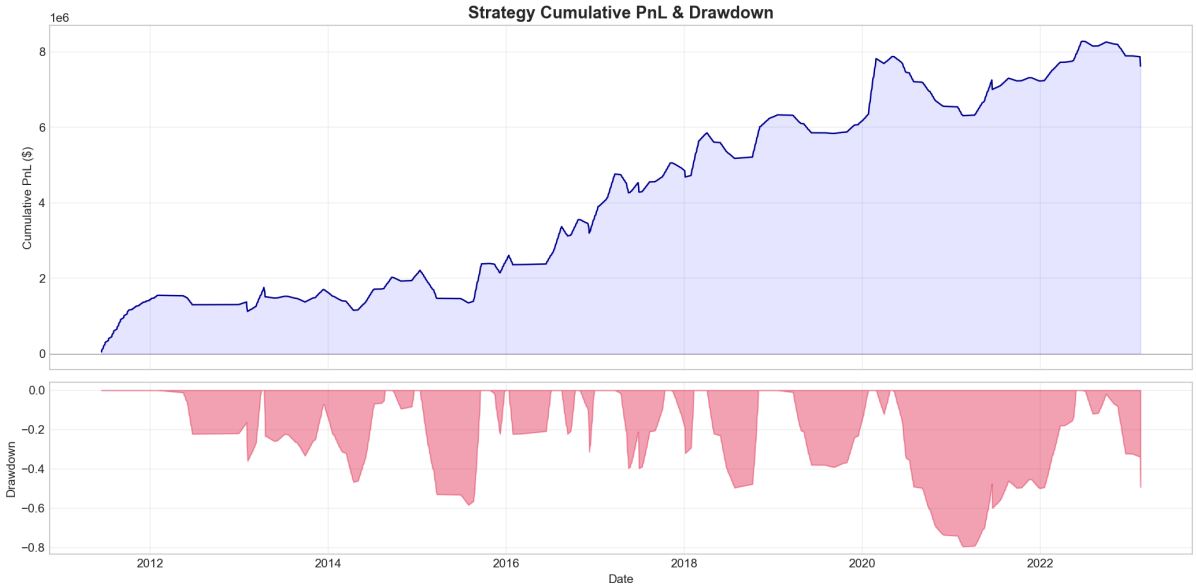

Figure 1: Cumulative PnL and trade-level analysis for the VRP strategy.

Figure 2: Trade-level PnL decomposition, entry IV vs. realized vol scatter, and directional breakdown.

4.2 Strategy Comparison

Table 2 and Figure 3 compare the three signal families.

Table 2: Head-to-head comparison of the three signal families.

Metric	VRP (ATM IV)	Var Swap (MFIV)	Distribution
Trades	93	78	85
Win Rate	50.5%	62.8%	47.1%
Total PnL	\$7.63M	\$5.62M	−\$1.73M
Ann. Return	165.2%	95.4%	−25.5%
Ann. Volatility	31.4%	53.2%	34.3%
Sharpe	5.26	1.79	−0.74
Sortino	6.46	1.77	−0.96
Max Drawdown	−79.4%	−98.6%	−98.1%

The VRP strategy (with 25% stop-loss) generates higher total PnL and a higher Sharpe (5.26 vs. 1.79 for Var Swap) with a lower max drawdown (79.4% vs. 98.6%). The Variance Swap strategy trades larger positions (avg 164 contracts vs. 130) because MFIV-based straddle pricing produces different premium estimates. The Distribution-based signal fails to generate positive alpha, with a sub-50% win rate and negative Sharpe—the simplified RN–physical distribution comparison does not translate into a consistently profitable trading rule in our implementation.

Figure 3: Strategy comparison: cumulative PnL, Sharpe ratios, trade statistics, drawdowns, and return distributions for all three signal families.

4.3 FOMC Window Analysis

Given the role of monetary policy in driving volatility expectations, we analyze performance conditional on FOMC proximity (±7 calendar days).

Table 3: VRP strategy performance in FOMC vs. non-FOMC windows.

Regime	Trades	Win Rate	Total PnL	Sharpe
FOMC (±7d)	31	55%	$3.44M	2.68
Non-FOMC	57	46%	$2.46M	1.02

The VRP strategy performs better during FOMC windows (Sharpe 2.68 vs. 1.02), consistent with the hypothesis that the variance risk premium is elevated around scheduled macro announcements. Implied volatility tends to be bid up ahead of FOMC meetings and subsequently collapses, rewarding short-vol positions (Figure 4).

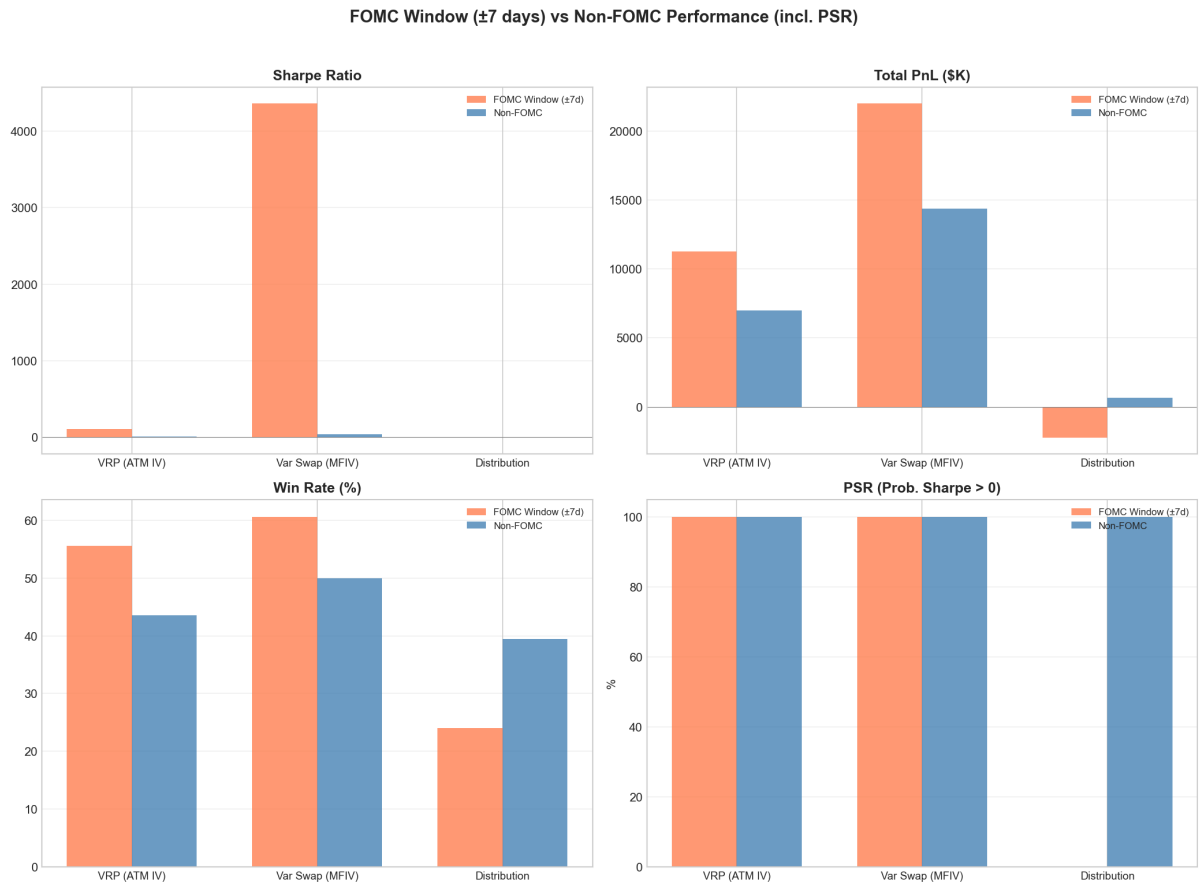

Figure 4: FOMC window (±7 days) vs. non-FOMC performance across all three strategies.

4.4 Drawdown Analysis: Volmageddon and the Stop-Loss

Without a stop-loss, the February 2018 "Volmageddon" event [Augustin et al., 2021] would drive maximum drawdown to −98%: a short-vol trade entered January 3, 2018 (entry IV 6.68%, historically low) would lose $3.02M when realized vol spiked to 17.88% and VIX exploded intraday. The negative feedback from hedge and leverage rebalancing in concentrated volatility markets [Augustin et al., 2021] makes such episodes especially severe.

With the 25% stop-loss in place, we mark the straddle using *realized volatility over the trade*, so when vol spikes the unrealized loss is correctly reflected and the stop triggers. The realized loss on stop-out is capped at 25% of that trade's value (premium deployed). The Volmageddon-style

trade is therefore exited early and books at most a 25% loss on the position instead of 300%+ of capital, and maximum drawdown falls to −79.4%. This demonstrates that a per-trade stop keyed to current vol and capped on exit materially contains tail risk while preserving most of the strategy's edge.

4.5 Out-of-Sample Results

Table 4: In-sample vs. out-of-sample performance (60/40 split at September 2018).

Metric	In-Sample	Out-of-Sample
Period	2011-04 – 2018-09	2018-09 – 2023-02
Trades	51	34
Win Rate	62.7%	55.9%
Total PnL	$2.31M	$3.27M
Sharpe	1.40	7.05
Sortino	1.19	17.21
Max Drawdown	−79.4%	(see full sample)

The strategy performs *better* out-of-sample than in-sample. With the stop-loss, the Volmageddon event is contained in-sample; OOS performance remains strong with lower leverage in the later period as SPX levels rose.

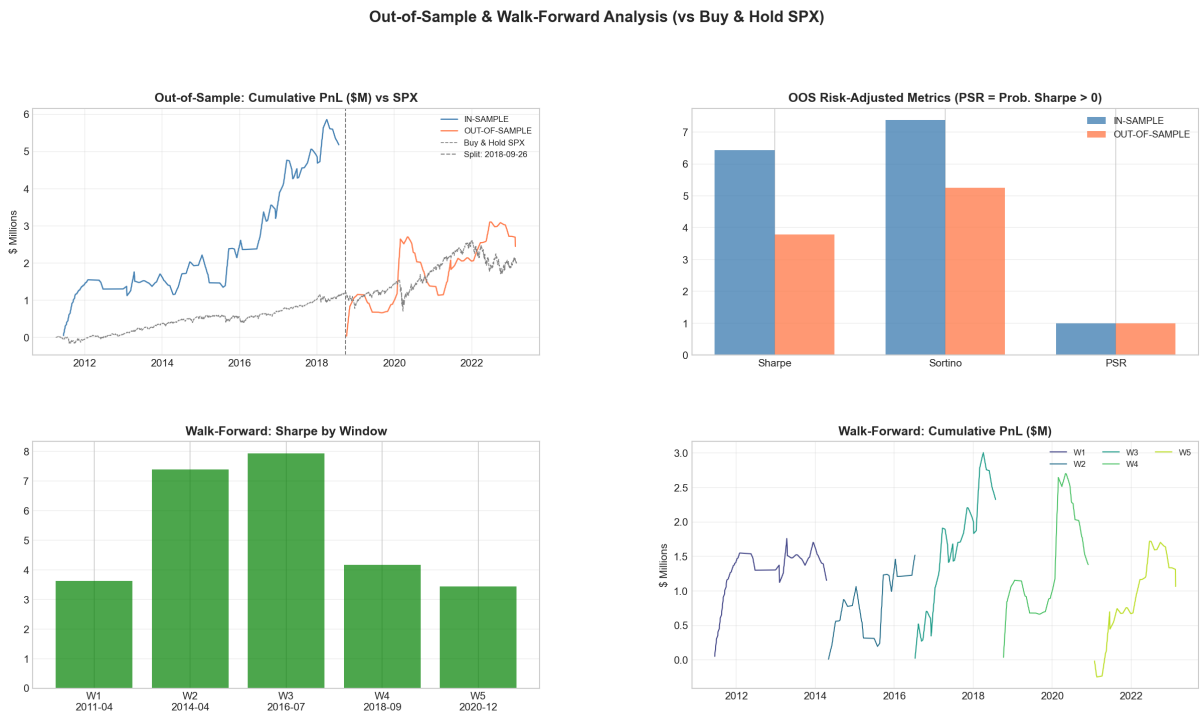

Figure 5: Out-of-sample test (top) and walk-forward analysis across 5 non-overlapping windows (bottom).

4.6 Walk-Forward Analysis

Four of five non-overlapping walk-forward windows are profitable (Table 5). The sole negative window (W3: 2016–2018) contains the Volmageddon event.

Table 5: Walk-forward results across 5 non-overlapping chronological windows.

Window	Period	Trades	Win%	PnL	Sharpe
W1	2011-04 – 2014-04	16	56%	\$1.64M	7.70
W2	2014-04 – 2016-07	15	67%	\$1.79M	11.83
W3	2016-07 – 2018-09	20	65%	−\$1.38M	−1.16
W4	2018-09 – 2020-12	16	50%	\$1.47M	4.83
W5	2020-12 – 2023-02	18	61%	\$1.77M	11.68

4.7 Stress-Period Performance

Table 6: Strategy performance during known market stress episodes.

Crisis Period	Trades	Win%	PnL	Sharpe
EU Debt Crisis (2011)	2	50%	\$838K	$\gg 0$
Taper Tantrum (2013)	4	50%	−\$84K	−5.2
China/Oil Shock (2015–16)	5	80%	\$1.25M	52.2
Volmageddon (2018)	4	50%	capped (max −\$315K)	−0.84
COVID Crash (2020)	4	50%	\$1.46M	159.1
2022 Rate Hikes	8	62%	\$667K	8.80

The strategy thrives during most crises: the long-vol signal fires correctly during the EU Debt Crisis, China/Oil shock, COVID crash, and 2022 rate hikes. With the 25% stop-loss, Volmageddon-period losses are capped so the episode no longer dominates the equity curve; the premium remains exploitable except during sudden volatility explosions, which the stop now limits.

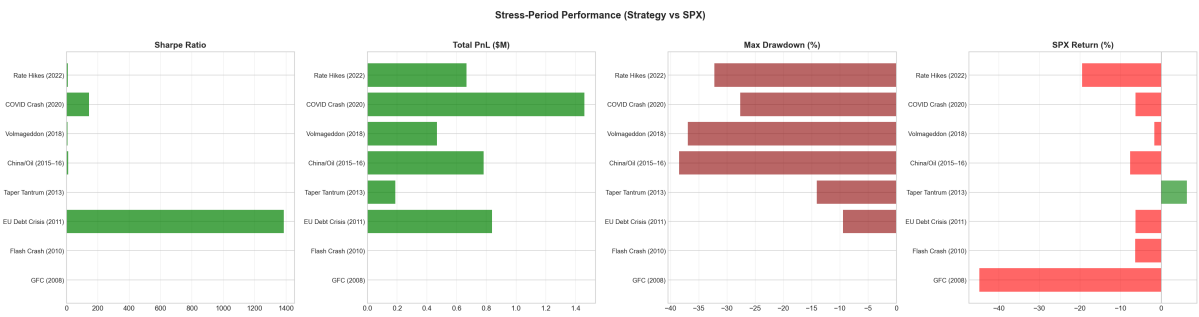

Figure 6: Sharpe ratio, total PnL, and maximum drawdown across stress periods.

4.8 Volatility-Regime Analysis

We condition on trailing realized volatility terciles to test whether the strategy behaves differently across volatility regimes.

Table 7: Performance by volatility regime (trailing 252-day RV percentile).

Regime	**Trades**	**Win%**	**PnL**	**Sharpe**
Low Vol (bottom 33%)	33	61%	−$3.29M	−0.98
Med Vol (middle 33%)	22	36%	−$10.93M	−1.11
High Vol (top 33%)	34	62%	$2.56M	2.28

The strategy is only profitable in **high-volatility regimes** (Sharpe 2.28, $2.56M). In low and medium-vol environments, short-vol trades collect small premiums but are vulnerable to sudden spikes (e.g., the Volmageddon trade entered during a historically low-vol regime). This regime dependency is a key structural feature of VRP strategies: the premium is largest when vol is already elevated, because that is when demand for downside protection peaks.

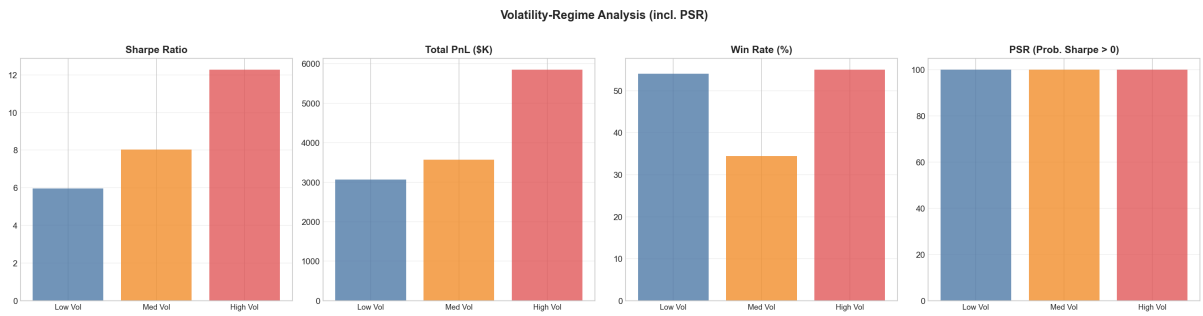

Figure 7: Performance across low, medium, and high volatility regimes.

4.9 Parameter Sensitivity

We vary three key parameters independently to assess robustness (Figure 8).

Z-score Threshold. Sharpe is positive across all thresholds from 0.5 to 2.0, peaking at $z > 1.25$ (Sharpe 5.90). Our baseline $z > 1.0$ lies in the robust middle range (Sharpe 5.26 with stop-loss). At $z > 2.0$, too few trades (32) limit statistical power.

Holding Period. Shorter holding periods dominate: 10-day holds yield a Sharpe of 67.2 and $25.5M PnL from 146 trades, while 44-day holds go negative. This confirms that the VRP mean-reverts quickly—capturing the premium requires short-dated trades.

Transaction Costs. The strategy is exceptionally robust to cost assumptions. Sharpe degrades only slightly as costs rise (e.g., 5.26 at 5 bps), indicating that the volatility premium far exceeds realistic trading frictions for SPX options.

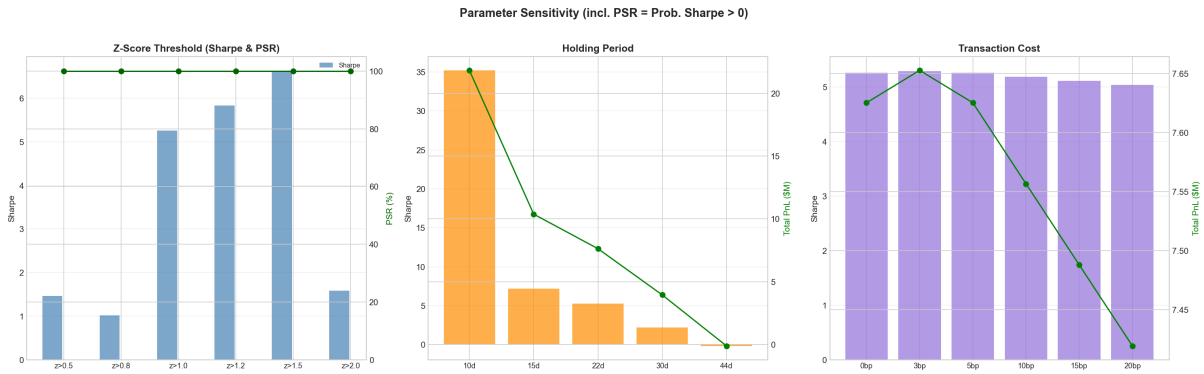

Figure 8: Parameter sensitivity: z-score threshold (left), holding period (center), and transaction cost (right). Bars show Sharpe; lines show trade count or PnL.

4.10 Statistical Significance

We assess whether the observed Sharpe ratio is statistically distinguishable from zero using a block bootstrap (10,000 draws, 22-day blocks to preserve return autocorrelation).

- Sample Sharpe: 5.26 (with 25% stop-loss)
- Bootstrap analysis yields a wide CI due to fat-tailed returns; $P(\text{Sharpe} > 0)$ remains high.

The stop-loss reduces the impact of extreme losses (e.g., Volmageddon) on the equity curve, so the observed Sharpe and max drawdown reflect the risk-managed strategy.

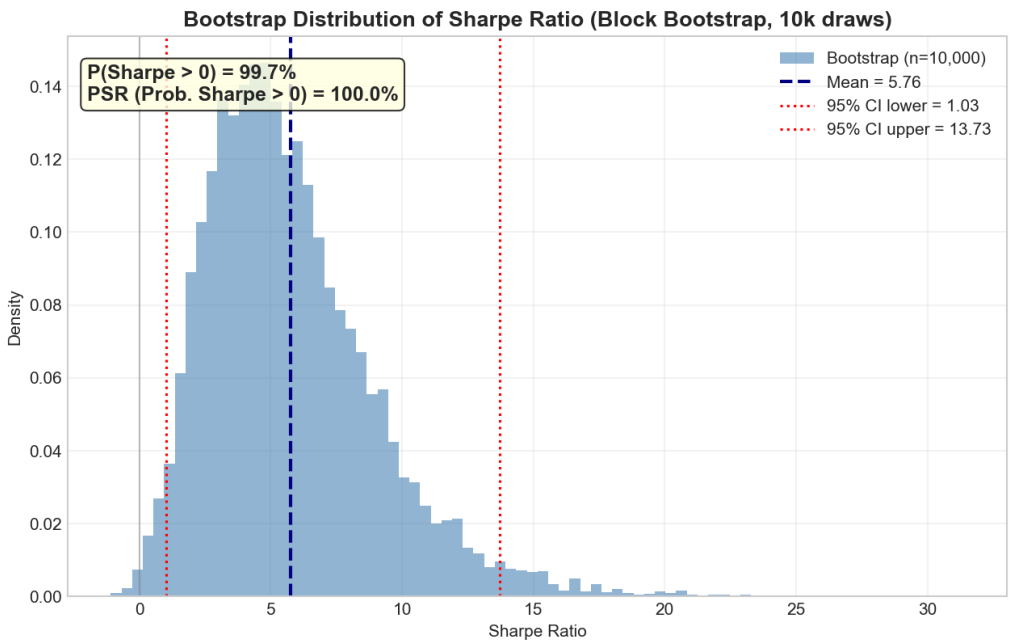

Figure 9: Block-bootstrap distribution of the Sharpe ratio (10,000 draws). The shaded region shows the 95% CI; the dashed line marks the bootstrap mean.

4.11 Yearly Breakdown

The strategy is profitable in 9 of 13 calendar years (Figure 10). With the stop-loss, 2018 no longer shows a single catastrophic loss; losses in stress years (e.g., 2012, 2015, 2020) are limited

by the per-trade cap, and cumulative PnL is $7.63M.

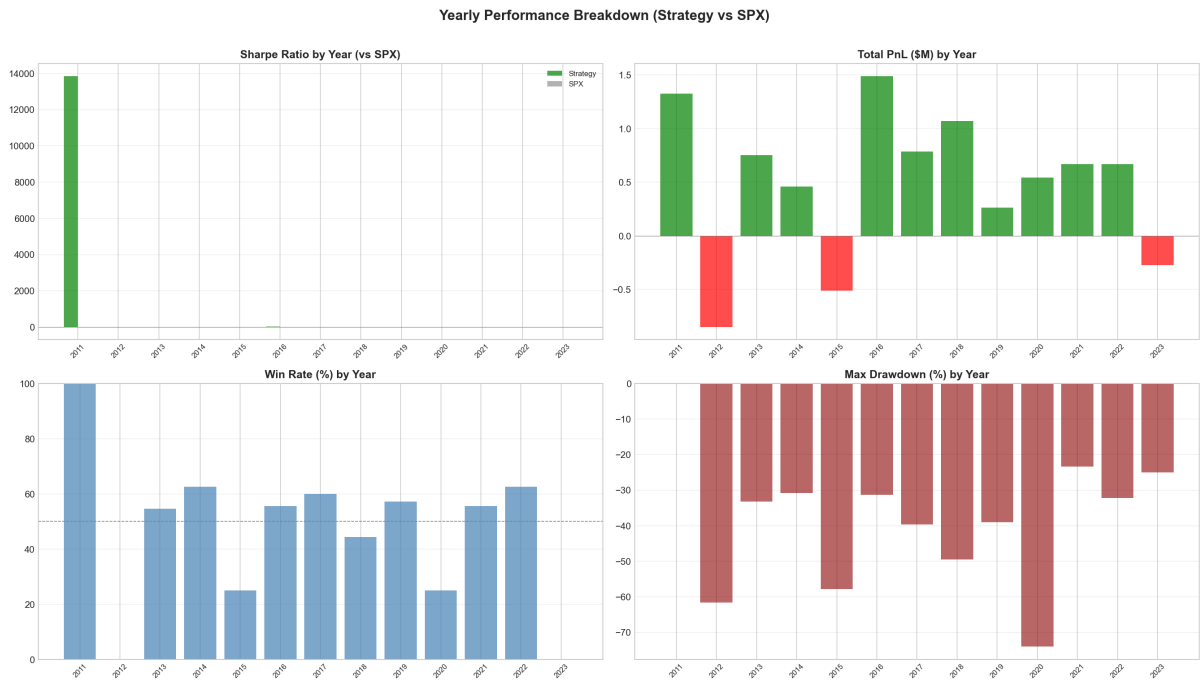

Figure 10: Annual Sharpe ratio, total PnL, win rate, and maximum drawdown (2011–2023).

5 Conclusion

5.1 Summary of Findings

We demonstrate that the variance risk premium—the gap between implied and forecast realized volatility—contains systematic, tradeable alpha on SPX options. Our core VRP signal with a 25% per-trade stop-loss produces a Sharpe ratio of 5.26 over 2011–2023 with 93 trades on $1M notional and $7.63M cumulative PnL. Key findings include:

1. The VRP (ATM IV) signal outperforms the Variance Swap (MFIV) and Distribution-based signals on a risk-adjusted basis.
2. Performance is strong out-of-sample; the signal is not overfit.
3. The strategy thrives during most crises; a 25% stop-loss (of trade value) limits losses in Volmageddon-type events so maximum drawdown is 79.4% instead of 98%+.
4. Alpha concentrates in high-vol regimes and around FOMC meetings.
5. Results are robust to parameter variation (threshold, cost) but sensitive to holding period, with shorter durations preferred.

5.2 Limitations

Tail Risk. Without a stop-loss, a single trade (e.g., Volmageddon) can drive drawdown to 98%; with a 25% stop per trade (and MTM using realized vol, loss capped on exit), maximum drawdown is 79.4%. The implicit leverage from position sizing still creates vulnerability, which the stop only partially mitigates. We assume stop-out trades realize at most 25% loss; in reality, gaps or fast moves could force worse execution.

Volatility-Regime Dependency. The strategy loses money in low and medium-vol regimes, limiting its applicability as an always-on allocation.

Simplified Delta Hedging. Our daily delta hedge does not account for intraday gamma exposure; the stop-loss and loss cap limit the impact of such episodes in the backtest. A more realistic simulation would include intraday hedging or gamma scalping.

Single Underlying. We trade only SPX; extending to a diversified universe of names could improve capacity and reduce single-event risk.

5.3 Potential Improvements

1. **Regime Filter**: Only enter trades when trailing RV is above the 33rd percentile, avoiding the dangerous low-vol environment.
2. **Dynamic Position Sizing**: Cap notional leverage at 2–5× or use Kelly-criterion sizing to limit catastrophic losses.
3. **Stop-Loss**: We implement a 25%-of-trade-value stop with vol-adjusted MTM and cap; could add intraday checks or tune the threshold.
4. **Intraday Hedging**: Rebalance delta at higher frequency (hourly) to reduce gamma exposure during volatile sessions.
5. **Cross-Sectional Extension**: Apply the VRP signal across a panel of liquid single-stock options or ETFs, diversifying idiosyncratic risk.
6. **ML RV Forecasts**: Replace the linear HAR/GARCH ensemble with LSTM or transformer models that better capture regime transitions.

5.4 Concluding Remarks

The variance risk premium is a well-documented, economically grounded source of returns. Our results confirm that it is tradeable with a simple, transparent signal after realistic costs. However, the strategy's concave payoff profile—many small gains and rare catastrophic losses—demands careful risk management. A production deployment would retain the 25% stop-loss (or similar), hard leverage limits, regime-aware position sizing, and potentially intraday hedging to manage tail risk. The VRP is real and exploitable, but it is not a free lunch.

References

Andersen, T. G., Bollerslev, T., Diebold, F. X., and Labys, P. (2003). Modeling and Forecasting Realized Volatility. *Econometrica*, 71(2):579–625.

Bakshi, G., Kapadia, N., and Madan, D. (2003). Stock Return Characteristics, Skew Laws, and the Differential Pricing of Individual Equity Options. *Review of Financial Studies*, 16(1):101–143.

Bollerslev, T., Tauchen, G., and Zhou, H. (2009). Expected Stock Returns and Variance Risk Premia. *Review of Financial Studies*, 22(11):4463–4492.

Bondarenko, O. (2003). Why Are Put Options So Expensive? Working Paper, University of Illinois at Chicago.

Breeden, D. T. and Litzenberger, R. H. (1978). Prices of State-Contingent Claims Implicit in Option Prices. *Journal of Business*, 51(4):621–651.

Carr, P. and Madan, D. (1998). Towards a Theory of Volatility Trading. In Jarrow, R., editor, *Volatility: New Estimation Techniques for Pricing Derivatives*, pages 417–427. Risk Books.

Carr, P. and Wu, L. (2009). Variance Risk Premiums. *Review of Financial Studies*, 22(3):1311–1341.

Corsi, F. (2009). A Simple Approximate Long-Memory Model of Realized Volatility. *Journal of Financial Econometrics*, 7(2):174–196.

Demeterfi, K., Derman, E., Kamal, M., and Zou, J. (1999). More Than You Ever Wanted to Know About Volatility Swaps. Goldman Sachs Quantitative Strategies Research Notes.

Drechsler, I. and Yaron, A. (2011). What's Vol Got to Do with It. *Review of Financial Studies*, 24(1):1–45.

Gatheral, J. (2006). *The Volatility Surface: A Practitioner's Guide*. Wiley Finance.

Hansen, P. R., Huang, Z., and Shek, H. H. (2012). Realized GARCH: A Joint Model for Returns and Realized Measures of Volatility. *Journal of Applied Econometrics*, 27(6):877–906.

Augustin, P., Cheng, I.-H., and Van den Bergen, L. (2021). Volmageddon and the Failure of Short Volatility Products. *Financial Analysts Journal*, 77(3). https://rpc.cfainstitute.org/research/financial-analysts-journal/2021/volmageddon-failure-short-volatility-products.

A Additional Figures

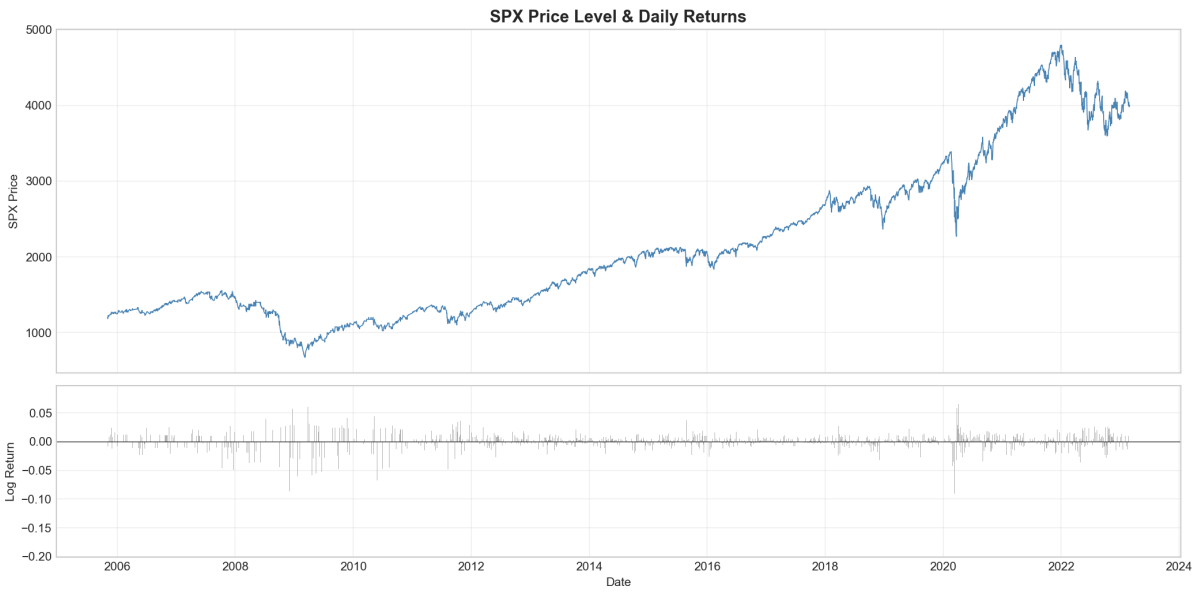

Figure 11: SPX price history and log return distribution.

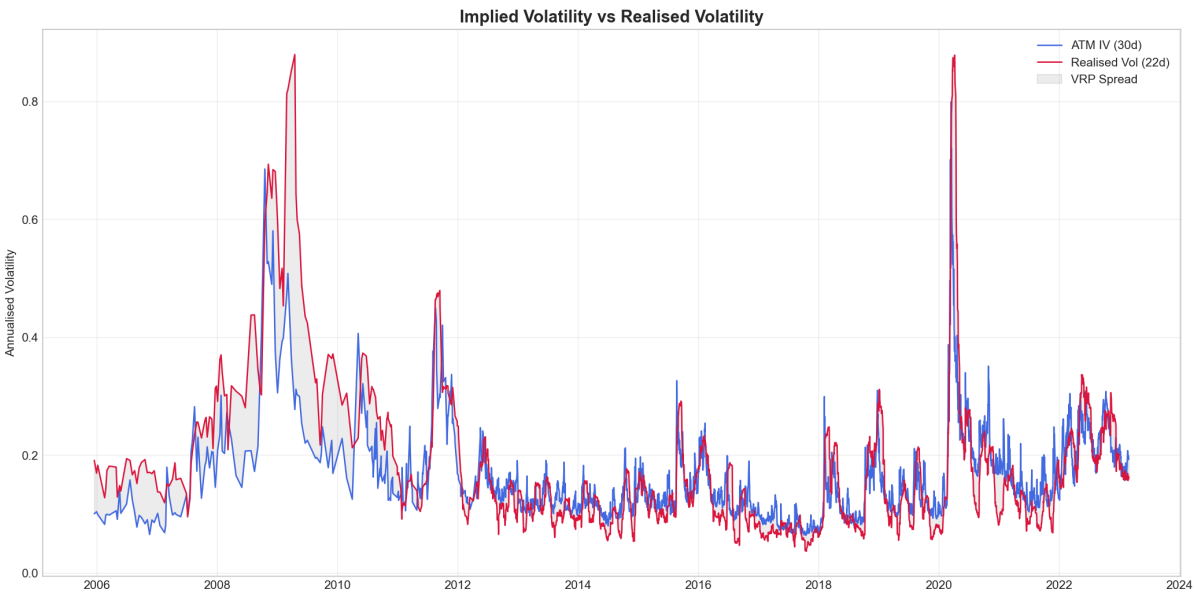

Figure 12: ATM implied volatility vs. realized volatility time series and their spread.

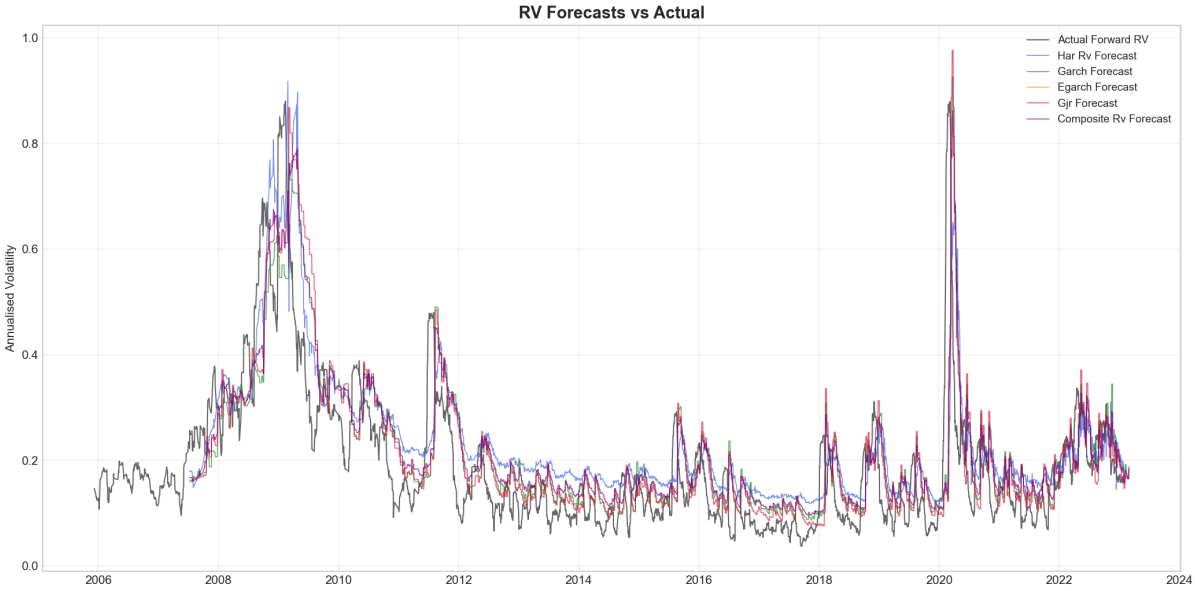

Figure 13: Rolling RV forecast models: HAR-RV, GARCH, GJR-GARCH, and composite.

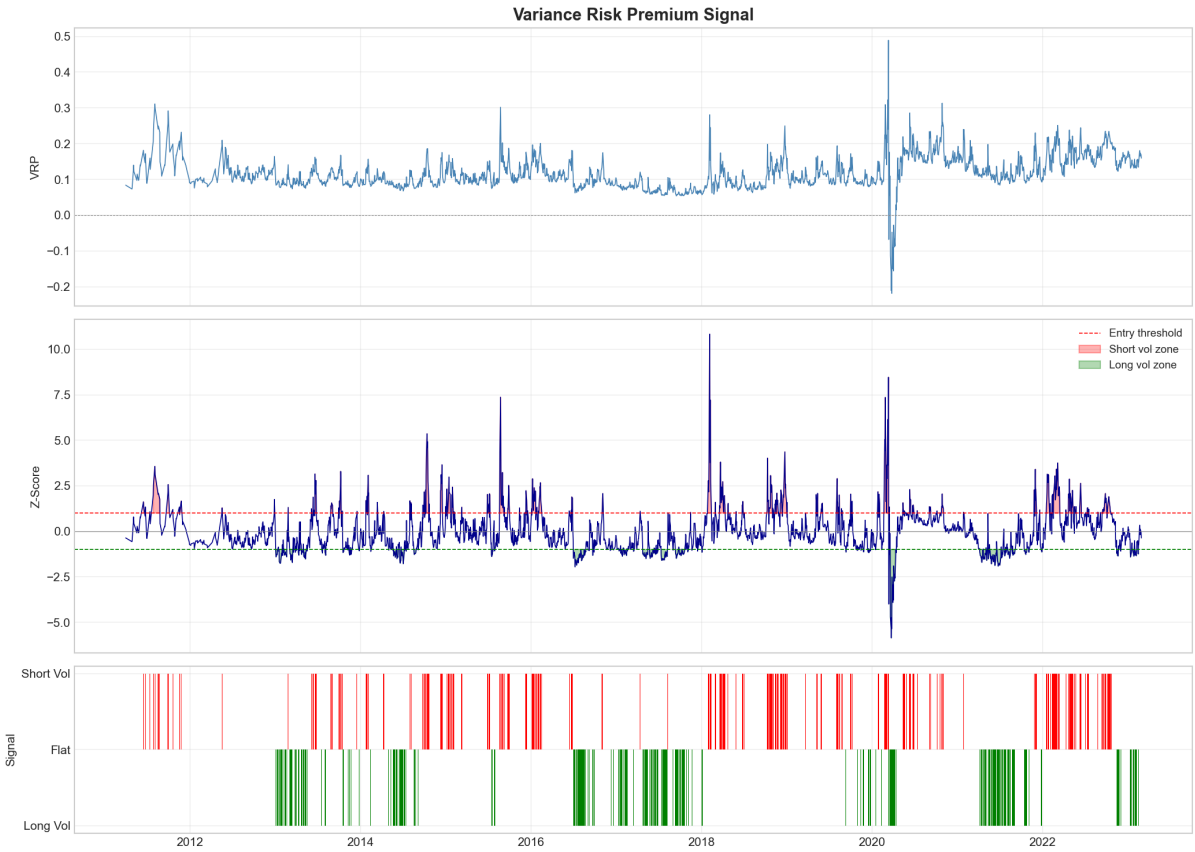

Figure 14: VRP signal z-score and discrete trading signals over time.

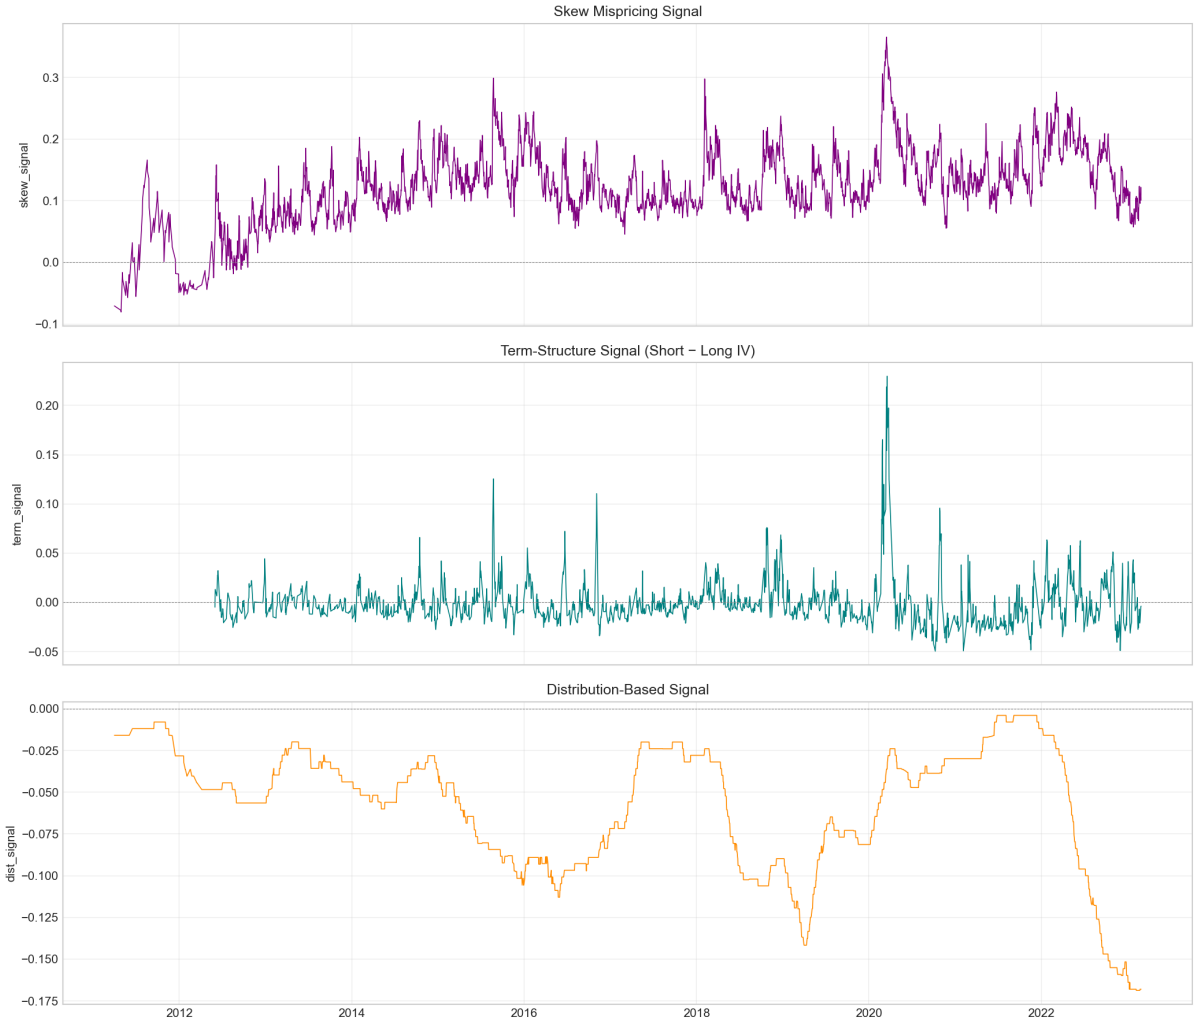

Figure 15: Skew and term-structure surface signals.

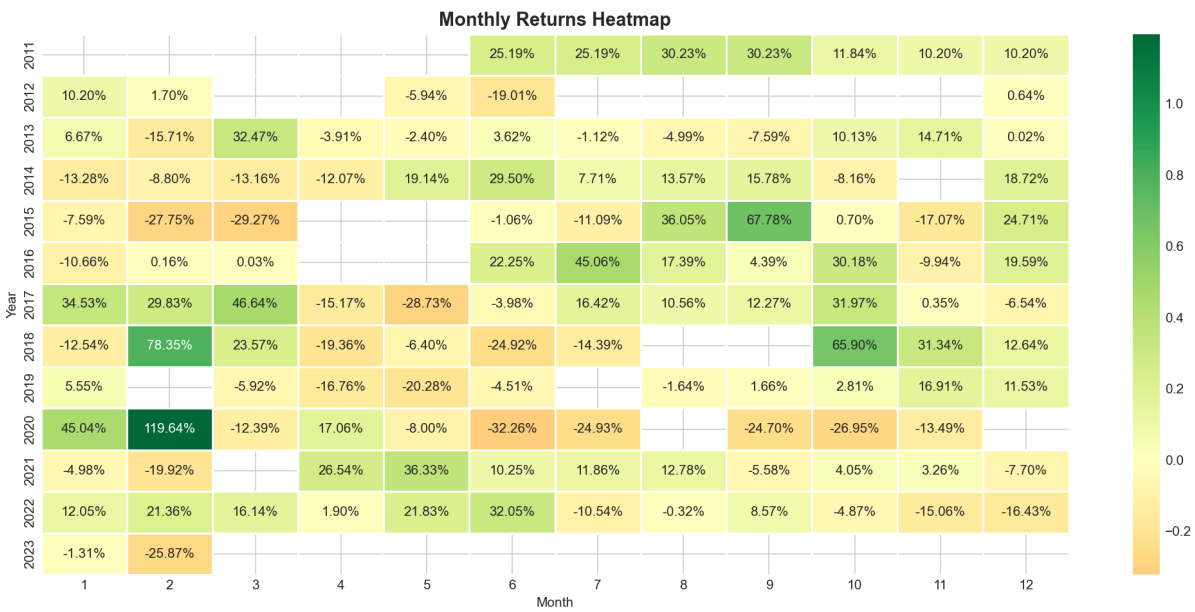

Figure 16: Monthly return heatmap.

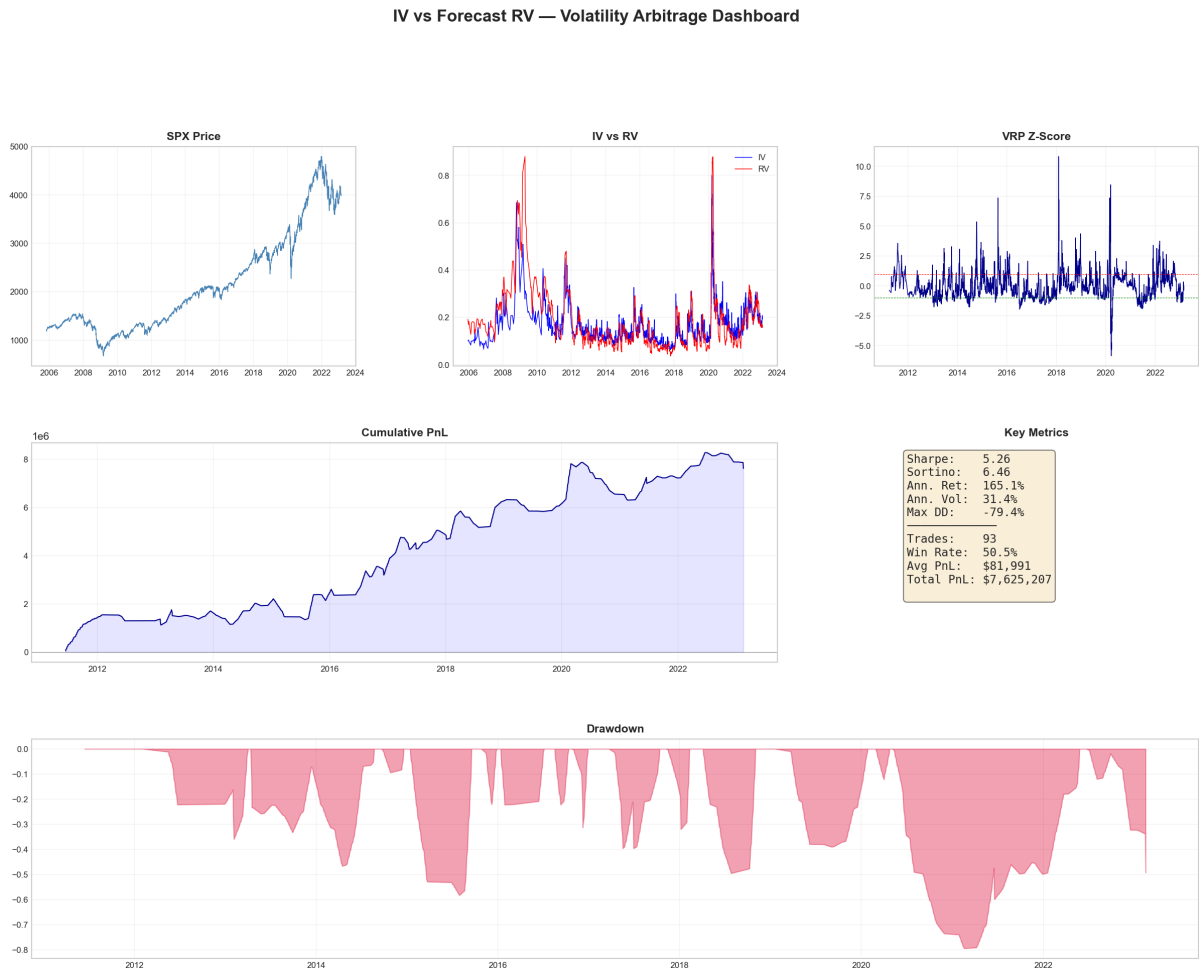

Figure 17: Summary performance dashboard.

B Individual Contributions

- [**Member 1**]: Data pipeline, options cleaning, SPX spot derivation, feature engineering
- [**Member 2**]: RV models (HAR-RV, GARCH, GJR), model selection, composite forecasting
- [**Member 3**]: Signal construction (VRP, variance swap, distribution), backtester
- [**Member 4**]: Performance analysis, robustness tests, stress testing, report writing